

# **Capstone Project-2**

## **(Employee Churn Prediction)**

### **Group-11**

**F-1295 Victor**

**F-1380 Habip**

**F-1470 Aziz**

**F-1491 Brian**

**F-1328 Fırat**

**F-1382 Halil**

# Table of Contents

- **Introduction to Employee Churn**
- **Key Takeaways and Assumptions**
- **EDA**
- **Data Visualization**
- **Data Pre-Processing**
- **Cluster Analysis**
- **Model Building**
- **Model Deployment**

# Introduction to Employee Churn

- In customer churn, you can predict who and when a customer will stop buying.
- Employee churn is **similar to customer churn**.
- It was found that employee churn will be affected by **age, tenure, pay, job satisfaction, salary, working conditions, growth potential and employee's perceptions of fairness**.
- Some other variables such as **gender, ethnicity, education, and marital status** were essential factors in the prediction of employee churn.
- In some cases such as the employee with **niche skills** are harder to replace.
- Acquiring new employees as a replacement has its costs such as **hiring costs and training costs**.
- Organizations tackle this problem by **applying machine learning techniques to predict employee churn**, which helps them in taking necessary actions.



# Introduction to Employee Churn

- **Business** chooses the employee to hire someone while in marketing **you don't get to choose your customers.**
- Employees will be the **face of your company**, and collectively, the employees produce everything your company does.
- **Losing a customer** affects **revenues and brand image.** Acquiring new customers is **difficult and costly** compared to retain the existing customer. **Employee churn** also **painful** for companies an organization. It requires **time and effort** in finding and training a replacement.
- **Employee churn** has **unique dynamics** compared to customer churn. It helps us in designing **better employee retention plans** and improving **employee satisfaction.** Data science algorithms can predict the future churn.



# Key Takeaways and Assumptions

- **A Clean DF (How lucky we are!🙏).**
- **There are no duplicated value.**
- **There are remarkable amount of niche employees.**
- **Clustering does not work for our project.**
- **XGBoost Model**
- **Model Deployment —> Streamlit**



# EDA

1.

Data Summary		Data Types	
dataframe	Values	Column Type	Count
Number of rows	14999	int32	6
Number of columns	10	float64	2
		string	2

skippy summary

column_name	NA	NA %	mean	sd	p0	p25	p75	p100	hist
satisfaction	0	0	0.61	0.25	0.09	0.44	0.82	1	
last_eval	0	0	0.72	0.17	0.36	0.56	0.87	1	
num_of_projects	0	0	3.8	1.2	2	3	5	7	
avg_monthly_hours	0	0	200	50	96	160	240	310	
experience	0	0	3.5	1.5	2	3	4	10	
work_accident	0	0	0.14	0.35	0	0	0	1	
promotion	0	0	0.021	0.14	0	0	0	1	
left	0	0	0.24	0.43	0	0	0	1	

number

column_name	NA	NA %	words per row	total words
department	0	0	1	15000
salary	0	0	1	15000

string

2.

We chose  
not to delete  
duplicated  
rows

3.

“groupby” function by  
[“work\_accident”,  
“promotion”,  
“department”,  
“salary”, “left”]

```
df.groupby("work_accident").mean() # promotion, left
```

	satisfaction	last_eval	num_of_projects	avg_monthly_hours	experience	promotion	left
work_accident							
0	0.607	0.717	3.805	201.259	3.497	0.019	0.265
1	0.648	0.713	3.789	199.818	3.506	0.035	0.078

Then, let's analyze employees according to their promotion status.

```
df.groupby("promotion").mean() # work accident, left
```

	satisfaction	last_eval	num_of_projects	avg_monthly_hours	experience	work_accident	left
promotion							
0	0.612	0.716	3.804	201.076	3.484	0.143	0.242
1	0.656	0.706	3.752	199.850	4.166	0.238	0.060

# EDA

## Analyze df column by column

### num\_of\_projects

```
In [237...] df.num_of_projects.value_counts()
```

```
Out[237...] 4    4365  
            3    4055  
            5    2761  
            2    2388  
            6    1174  
            7     256  
Name: num_of_projects, dtype: int64
```

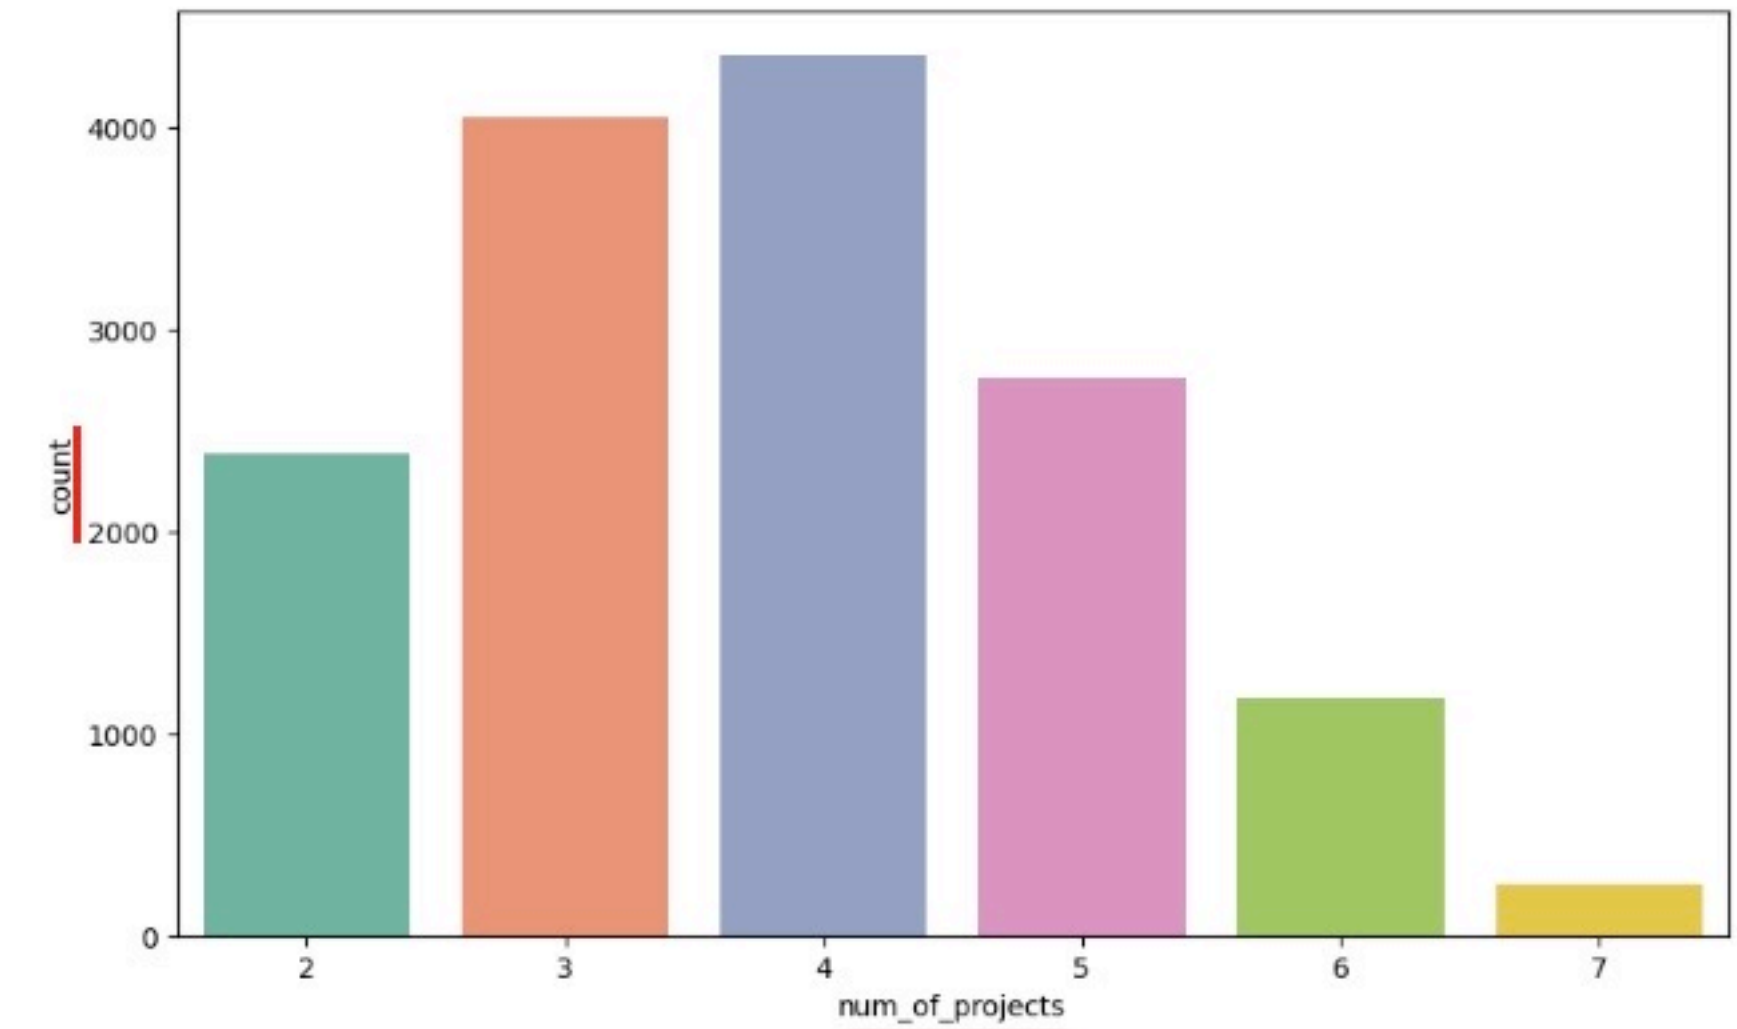
```
In [238...] df.num_of_projects.value_counts(normalize=True)
```

```
Out[238...] 4    0.291  
            3    0.270  
            5    0.184  
            2    0.159  
            6    0.078  
            7    0.017  
Name: num_of_projects, dtype: float64
```


```
In [239...] df.num_of_projects.describe()
```

```
Out[239...] count    14999.000  
            mean         3.803  
            std         1.233  
            min         2.000  
            25%         3.000  
            50%         4.000  
            75%         5.000  
            max         7.000  
Name: num_of_projects, dtype: float64
```

```
0... sns.countplot(x=df.num_of_projects, palette="Set2");
```




```
1... sns.countplot(x="num_of_projects", hue="left", palette="Set2", data=df);
```



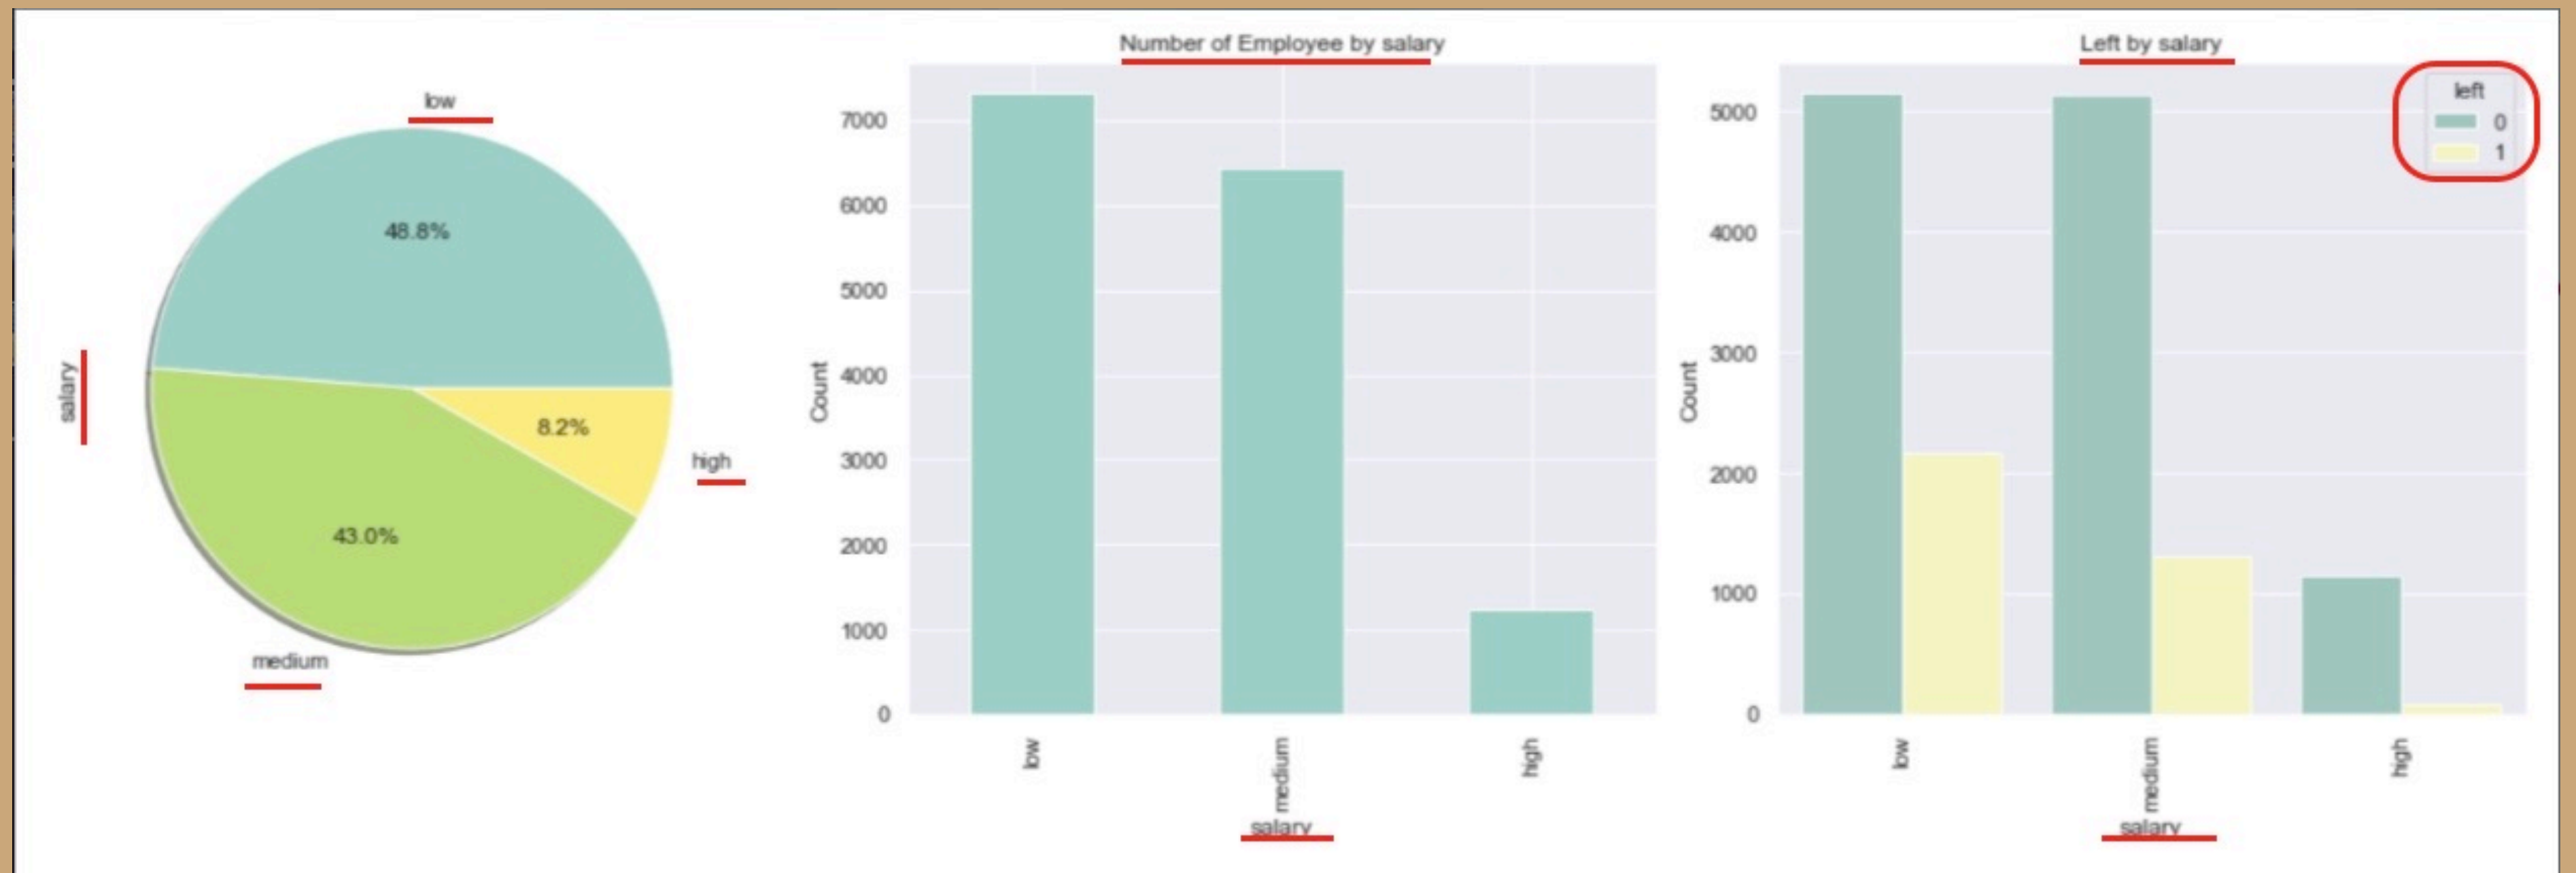


# Data Visualization

1.



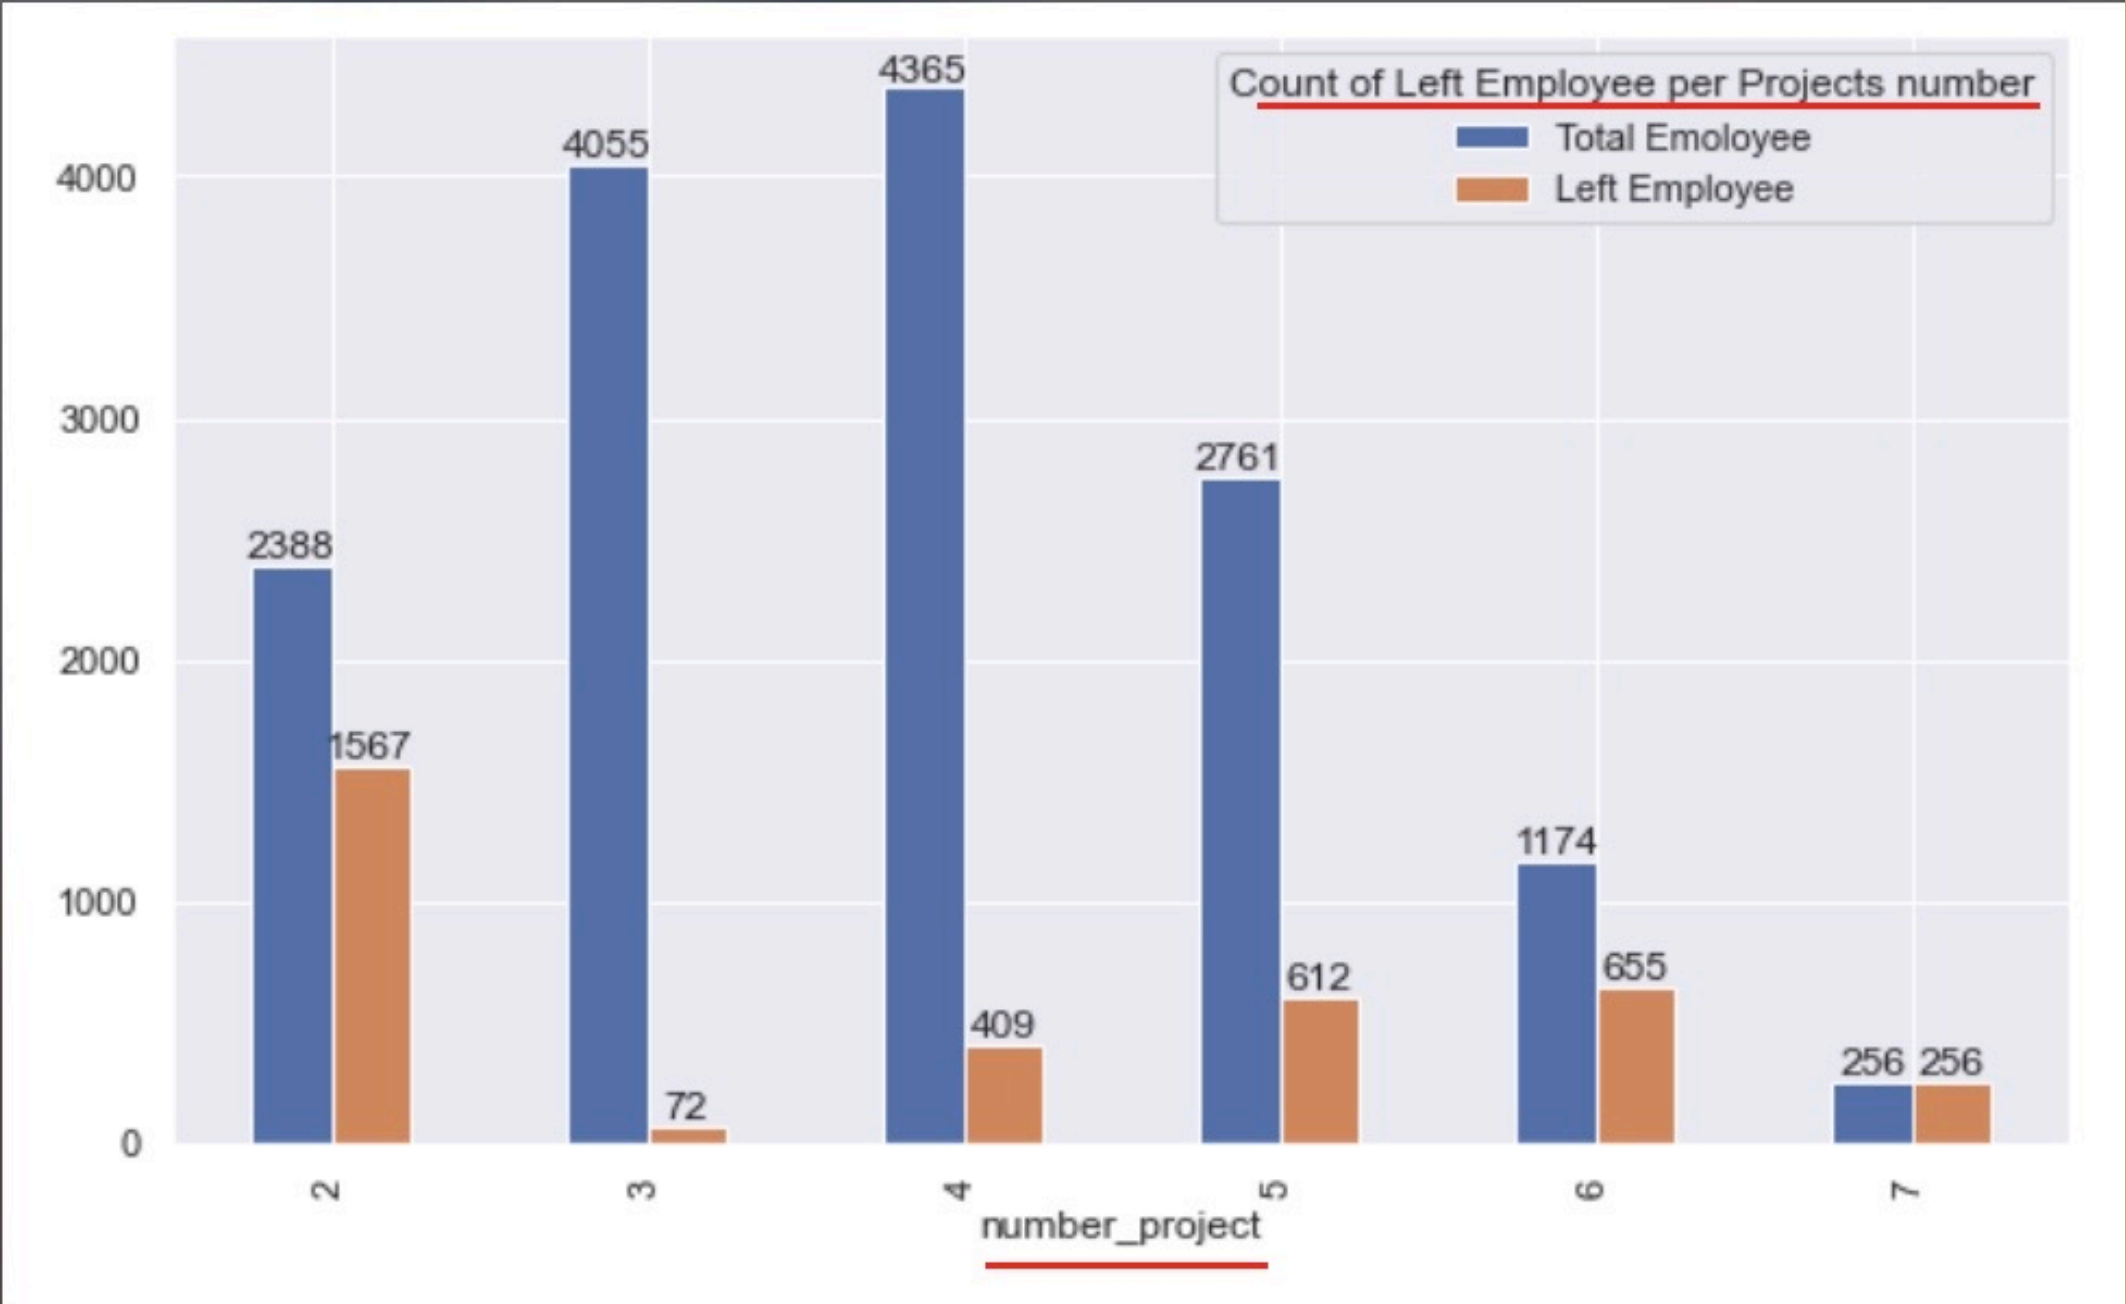
2.



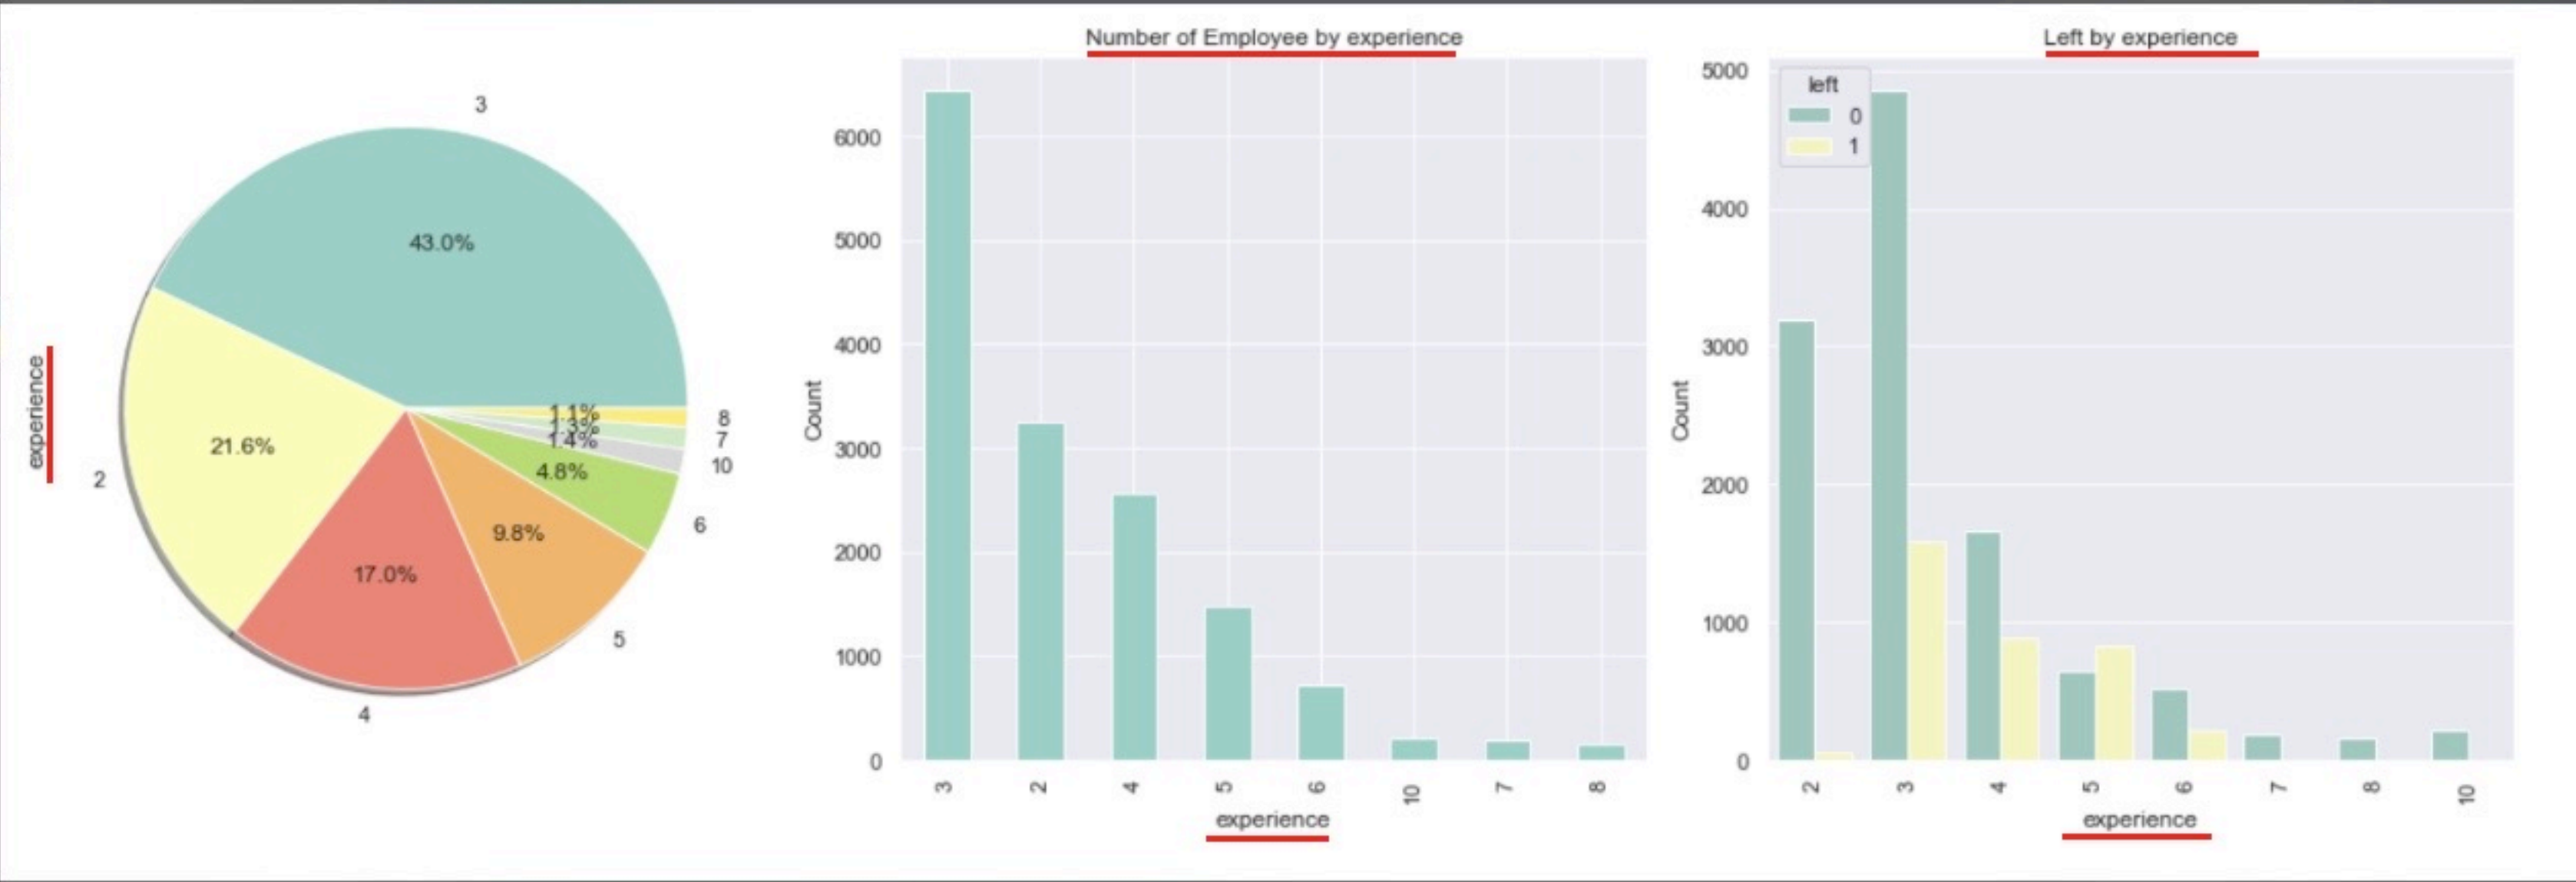


# Data Visualization


1.



2.



# Data Visualization



# Data Pre-Processing

```
▼ column_trans = make_column_transformer(  
    (OneHotEncoder(handle_unknown="ignore", sparse=False), ["department"]),  
    (OrdinalEncoder(handle_unknown='use_encoded_value', unknown_value=-1), ["salary"]),  
    remainder=MinMaxScaler()  
)  
X = pd.DataFrame(data=column_trans.fit_transform(X), columns=column_trans.get_feature_names_out())
```

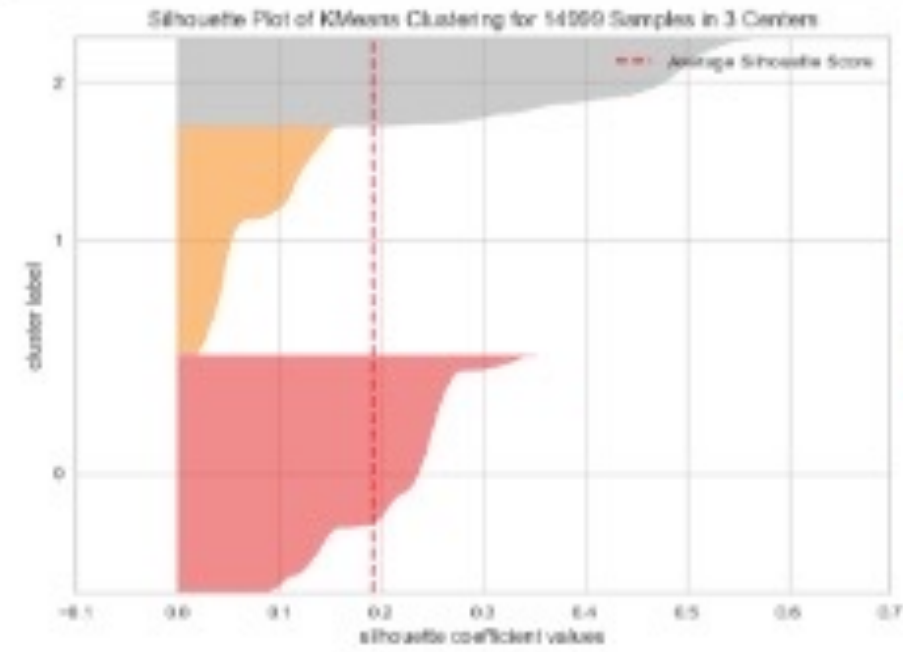
**OneHotEncoder [“department”]**

**OrdinalEncoder [“salary”]**

**MinMaxScaler —> Remainder**



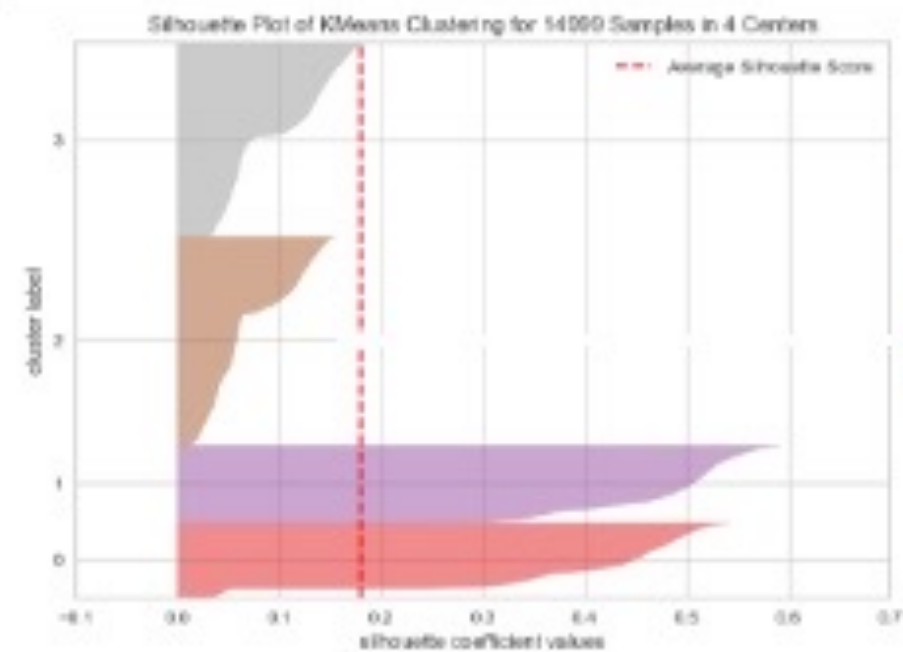
# Cluster Analysis



3

```
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score

model_1 = KMeans(n_clusters = 3, random_state = 100)
silhouette = silhouette_score(model_1)
silhouette_plot = silhouette_plot(model_1)
silhouette_plot.show()
```



4

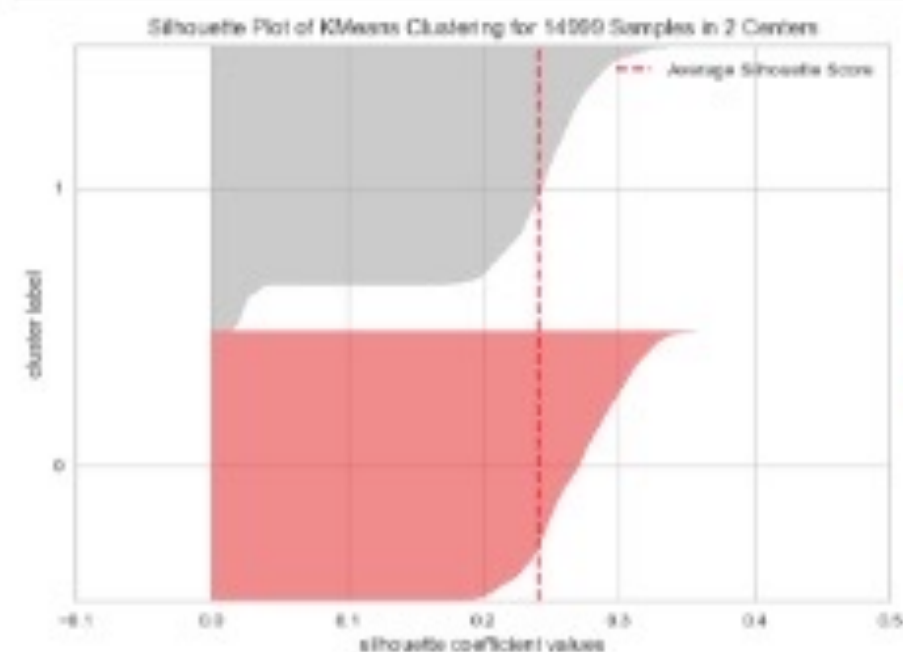
According to the silhouette score, clustering according to the K=2, K=3 and K=4 are seen above.

For K=2, 0 labelled cluster is below the average silhouette score. For K=3, 0 and 3 labelled clusters are below the average

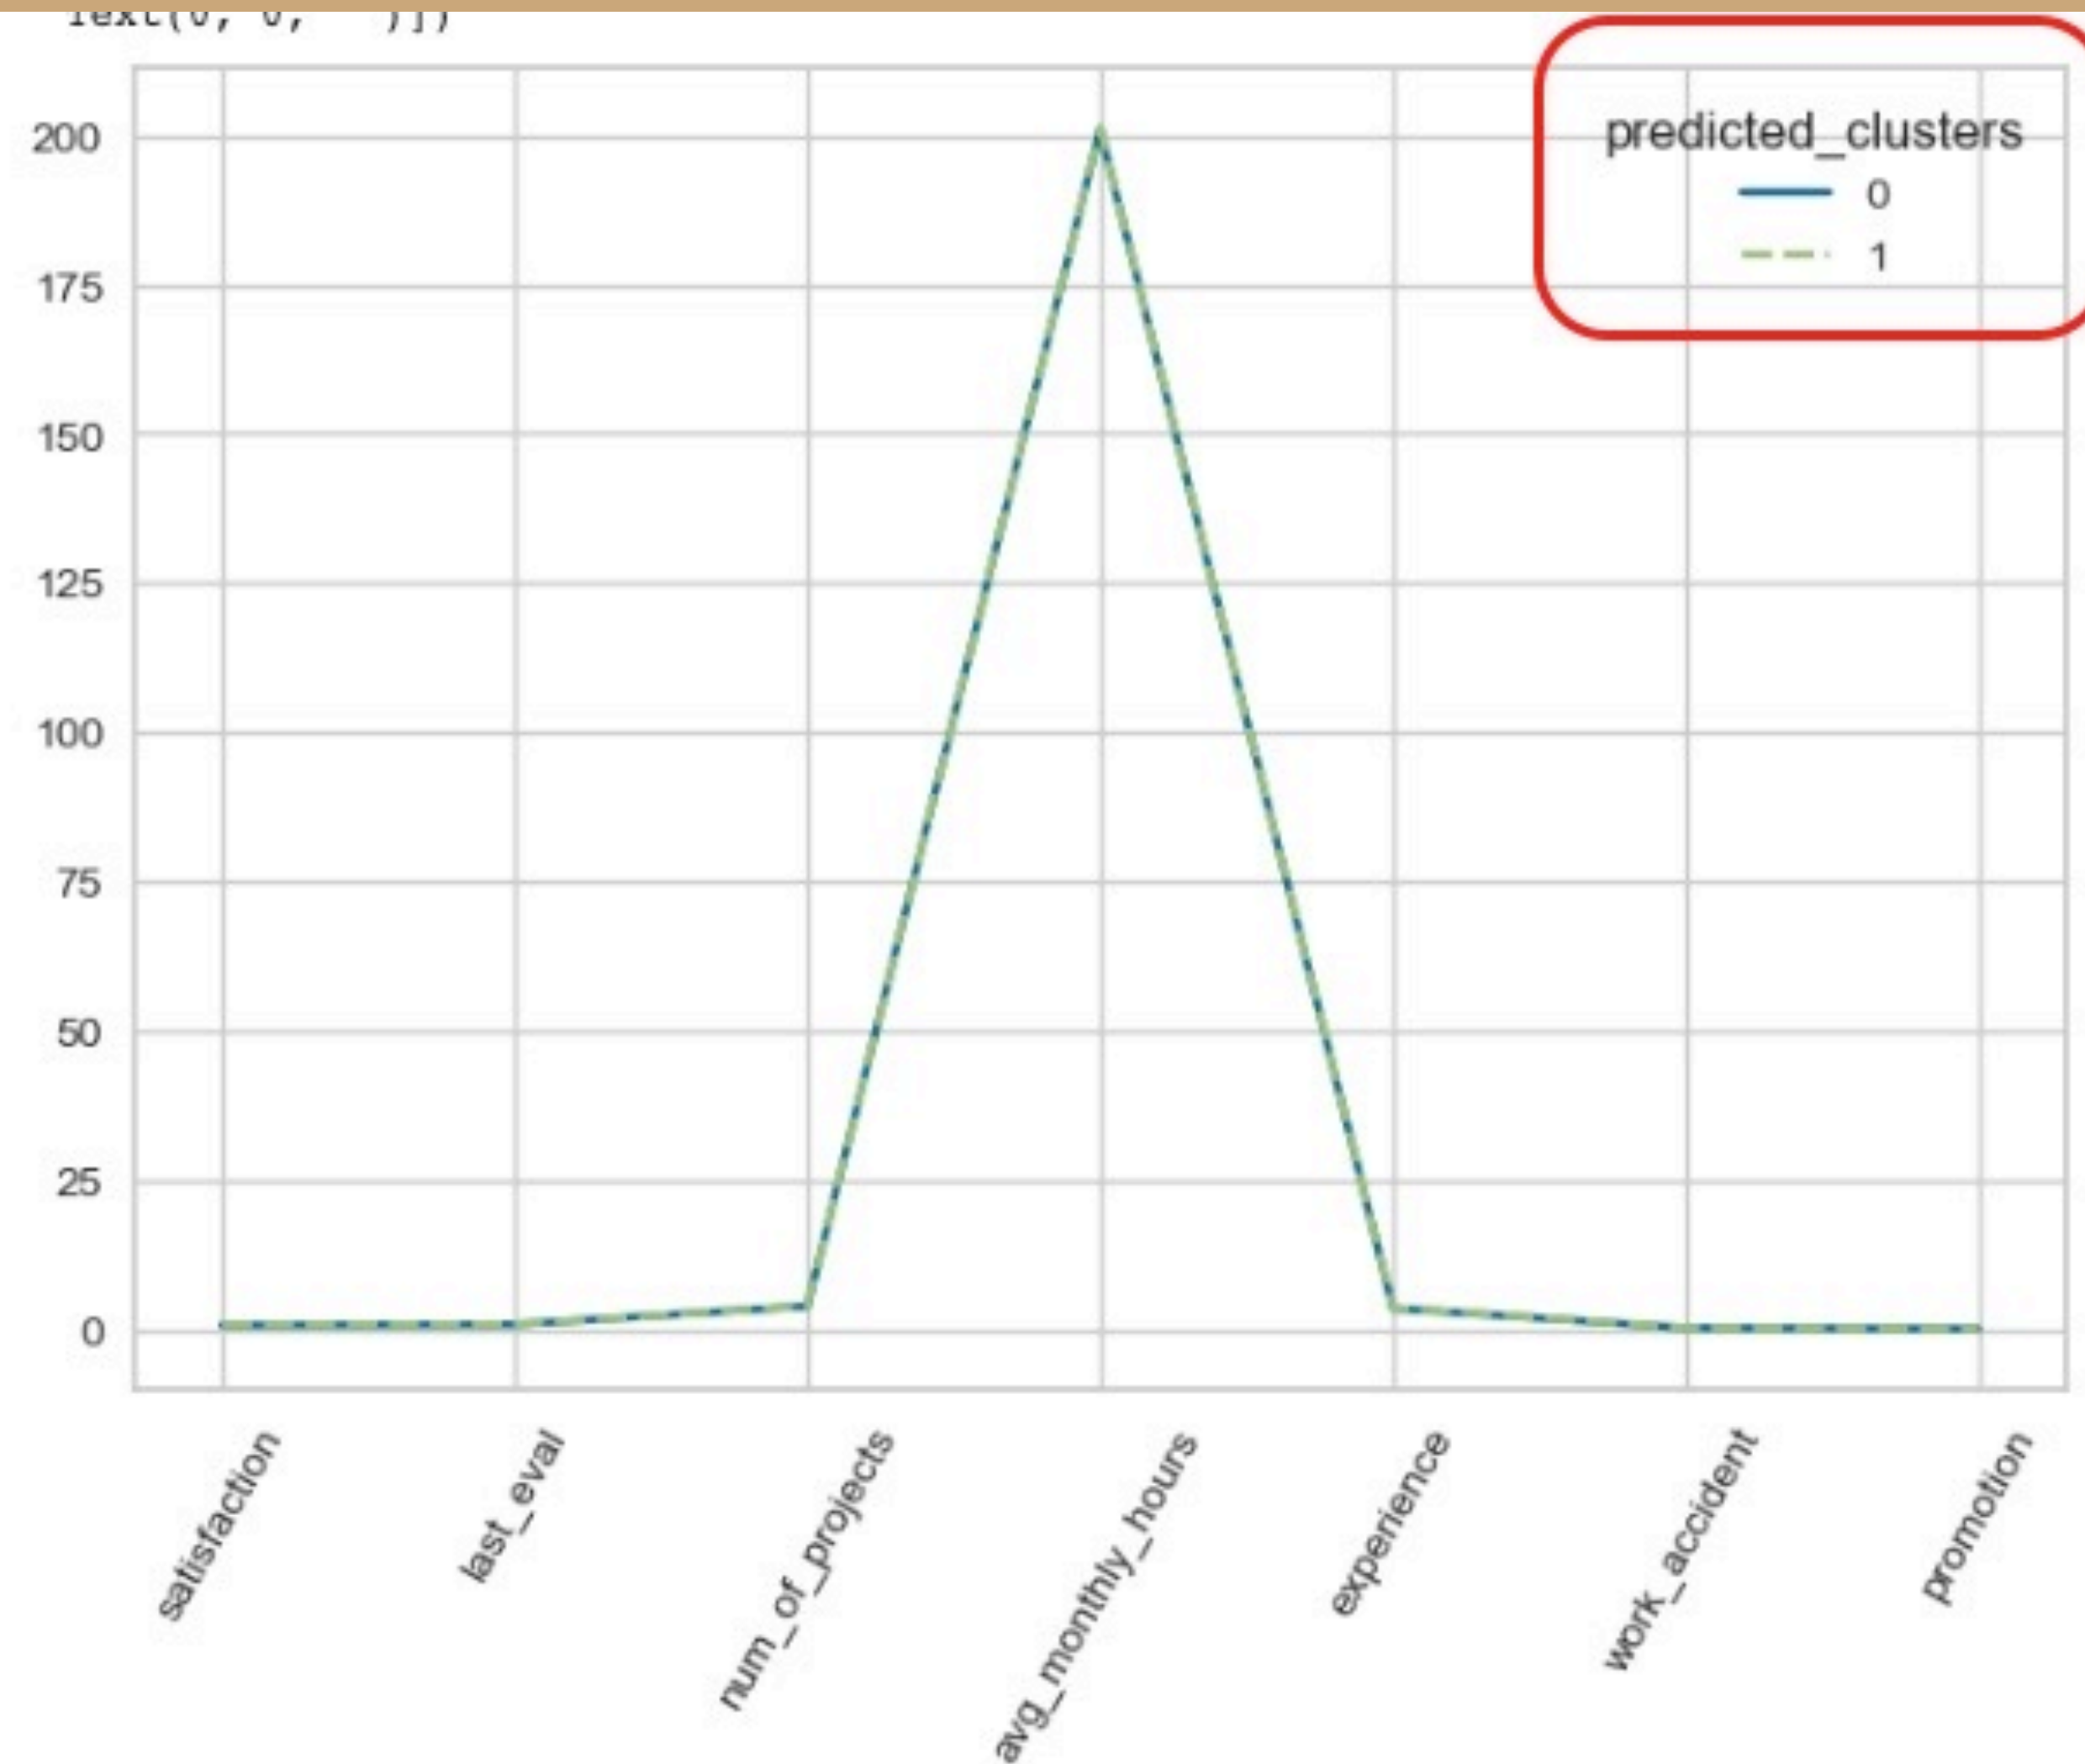
Let's see how it is when K=2 (According to our target variable classes)

```
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score

model_2 = KMeans(n_clusters = 2, random_state = 100)
silhouette = silhouette_score(model_2)
silhouette_plot = silhouette_plot(model_2)
silhouette_plot.show()
```



2



As seen above, the columns in the data set do not separate from each other. All columns are intertwined with each other.

As seen above it is visually obvious that clustering is not a good approach to our dataset.

From now on we are going to use classification models to make churn predictions.



# Model Building

## Logistic Regression

```
[565... eval_metric(grid_model, X_train, y_train, X_test, y_test)
```

### Test Set

```
[[873 270]
 [ 75 282]]
```


	precision	<u>recall</u>	f1-score	support
0	0.92	0.76	0.84	1143
<u>1</u>	0.51	<u>0.79</u>	0.62	357
<u>accuracy</u>			<u>0.77</u>	1500
macro avg	0.72	0.78	0.73	1500
weighted avg	0.82	0.77	0.78	1500

### Train Set

```
[[7616 2669]
 [ 693 2521]]
```

	precision	<u>recall</u>	f1-score	support
0	0.92	0.74	0.82	10285
<u>1</u>	0.49	<u>0.78</u>	0.60	3214
<u>accuracy</u>			<u>0.75</u>	13499
macro avg	0.70	0.76	0.71	13499
weighted avg	0.81	0.75	0.77	13499

```
[551... plot_confusion_matrix(log_model, X_test, y_test);
```



# Model Building

## KNN

### Test\_Set

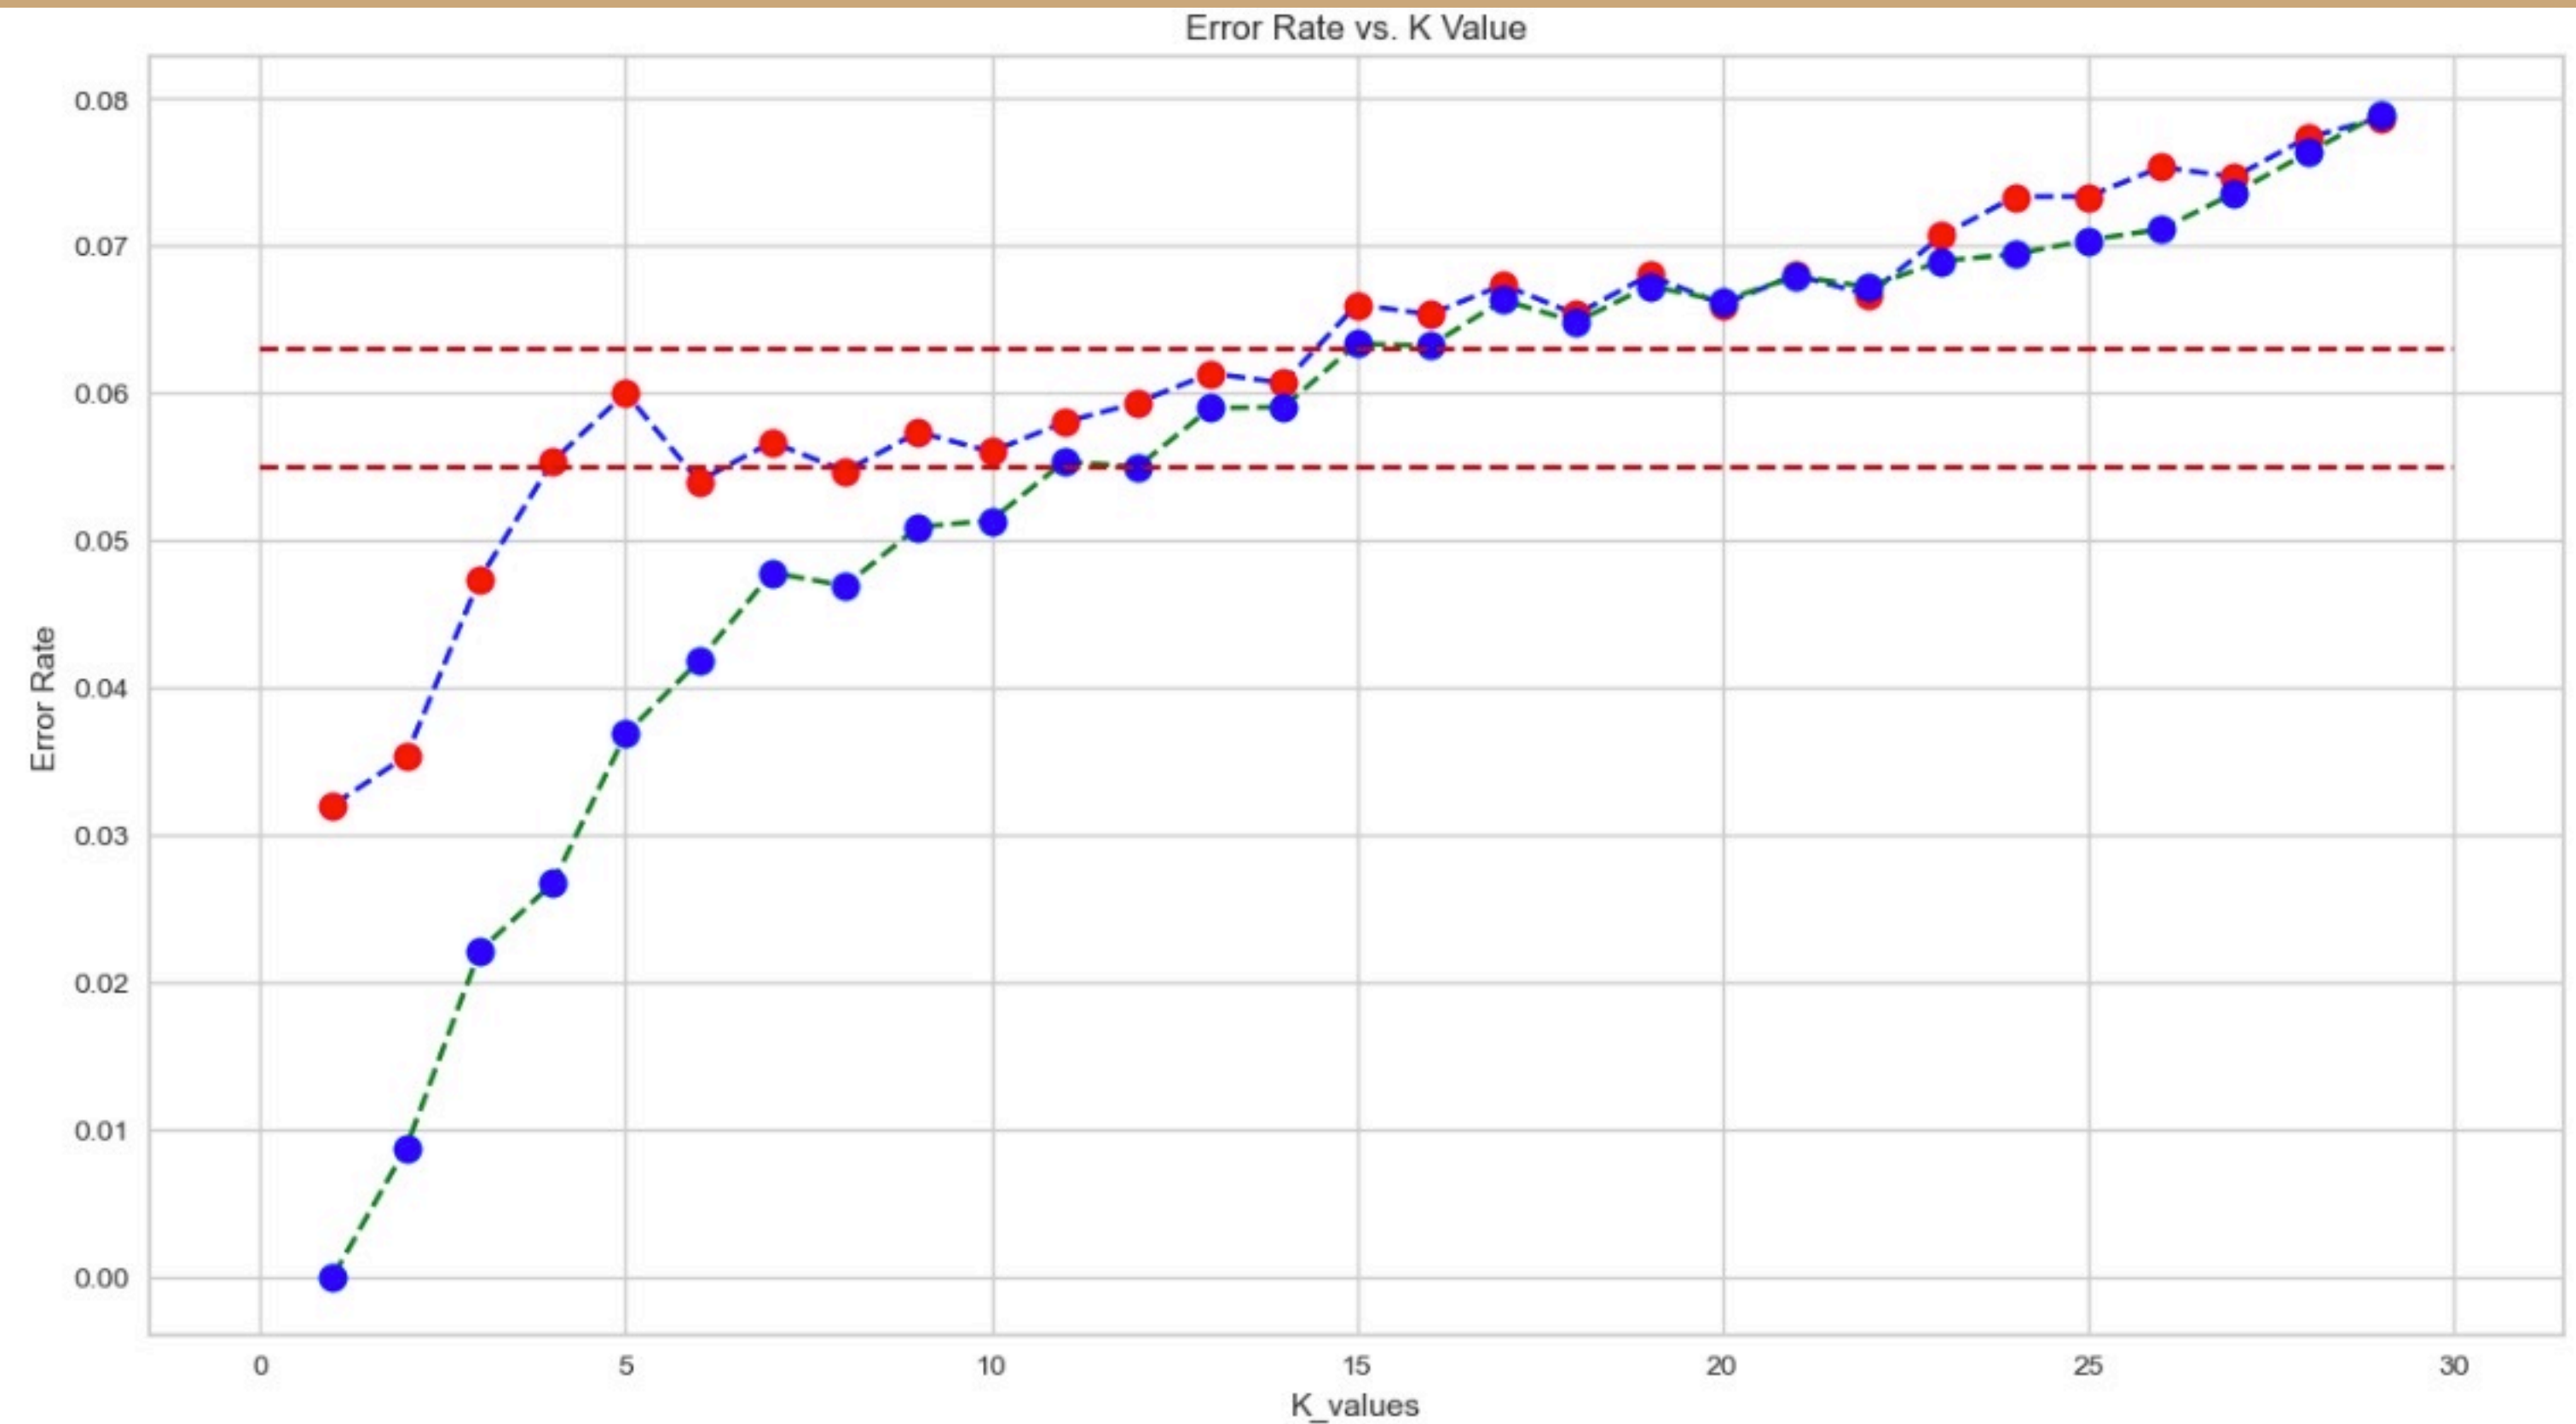
```
[[1101  42]
 [  15 342]]
```

	precision	<u>recall</u>	f1-score	support
0	0.99	0.96	0.97	1143
<u>1</u>	0.89	<u>0.96</u>	0.92	357
<u>accuracy</u>			<u>0.96</u>	1500
macro avg	0.94	0.96	0.95	1500
weighted avg	0.96	0.96	0.96	1500

### Train\_Set

```
[[10285  0]
 [   0 3214]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10285
<u>1</u>	1.00	<u>1.00</u>	1.00	3214
<u>accuracy</u>			<u>1.00</u>	13499
macro avg	1.00	1.00	1.00	13499
weighted avg	1.00	1.00	1.00	13499



# Model Building

## RainForest

RF\_tuned

Test Set

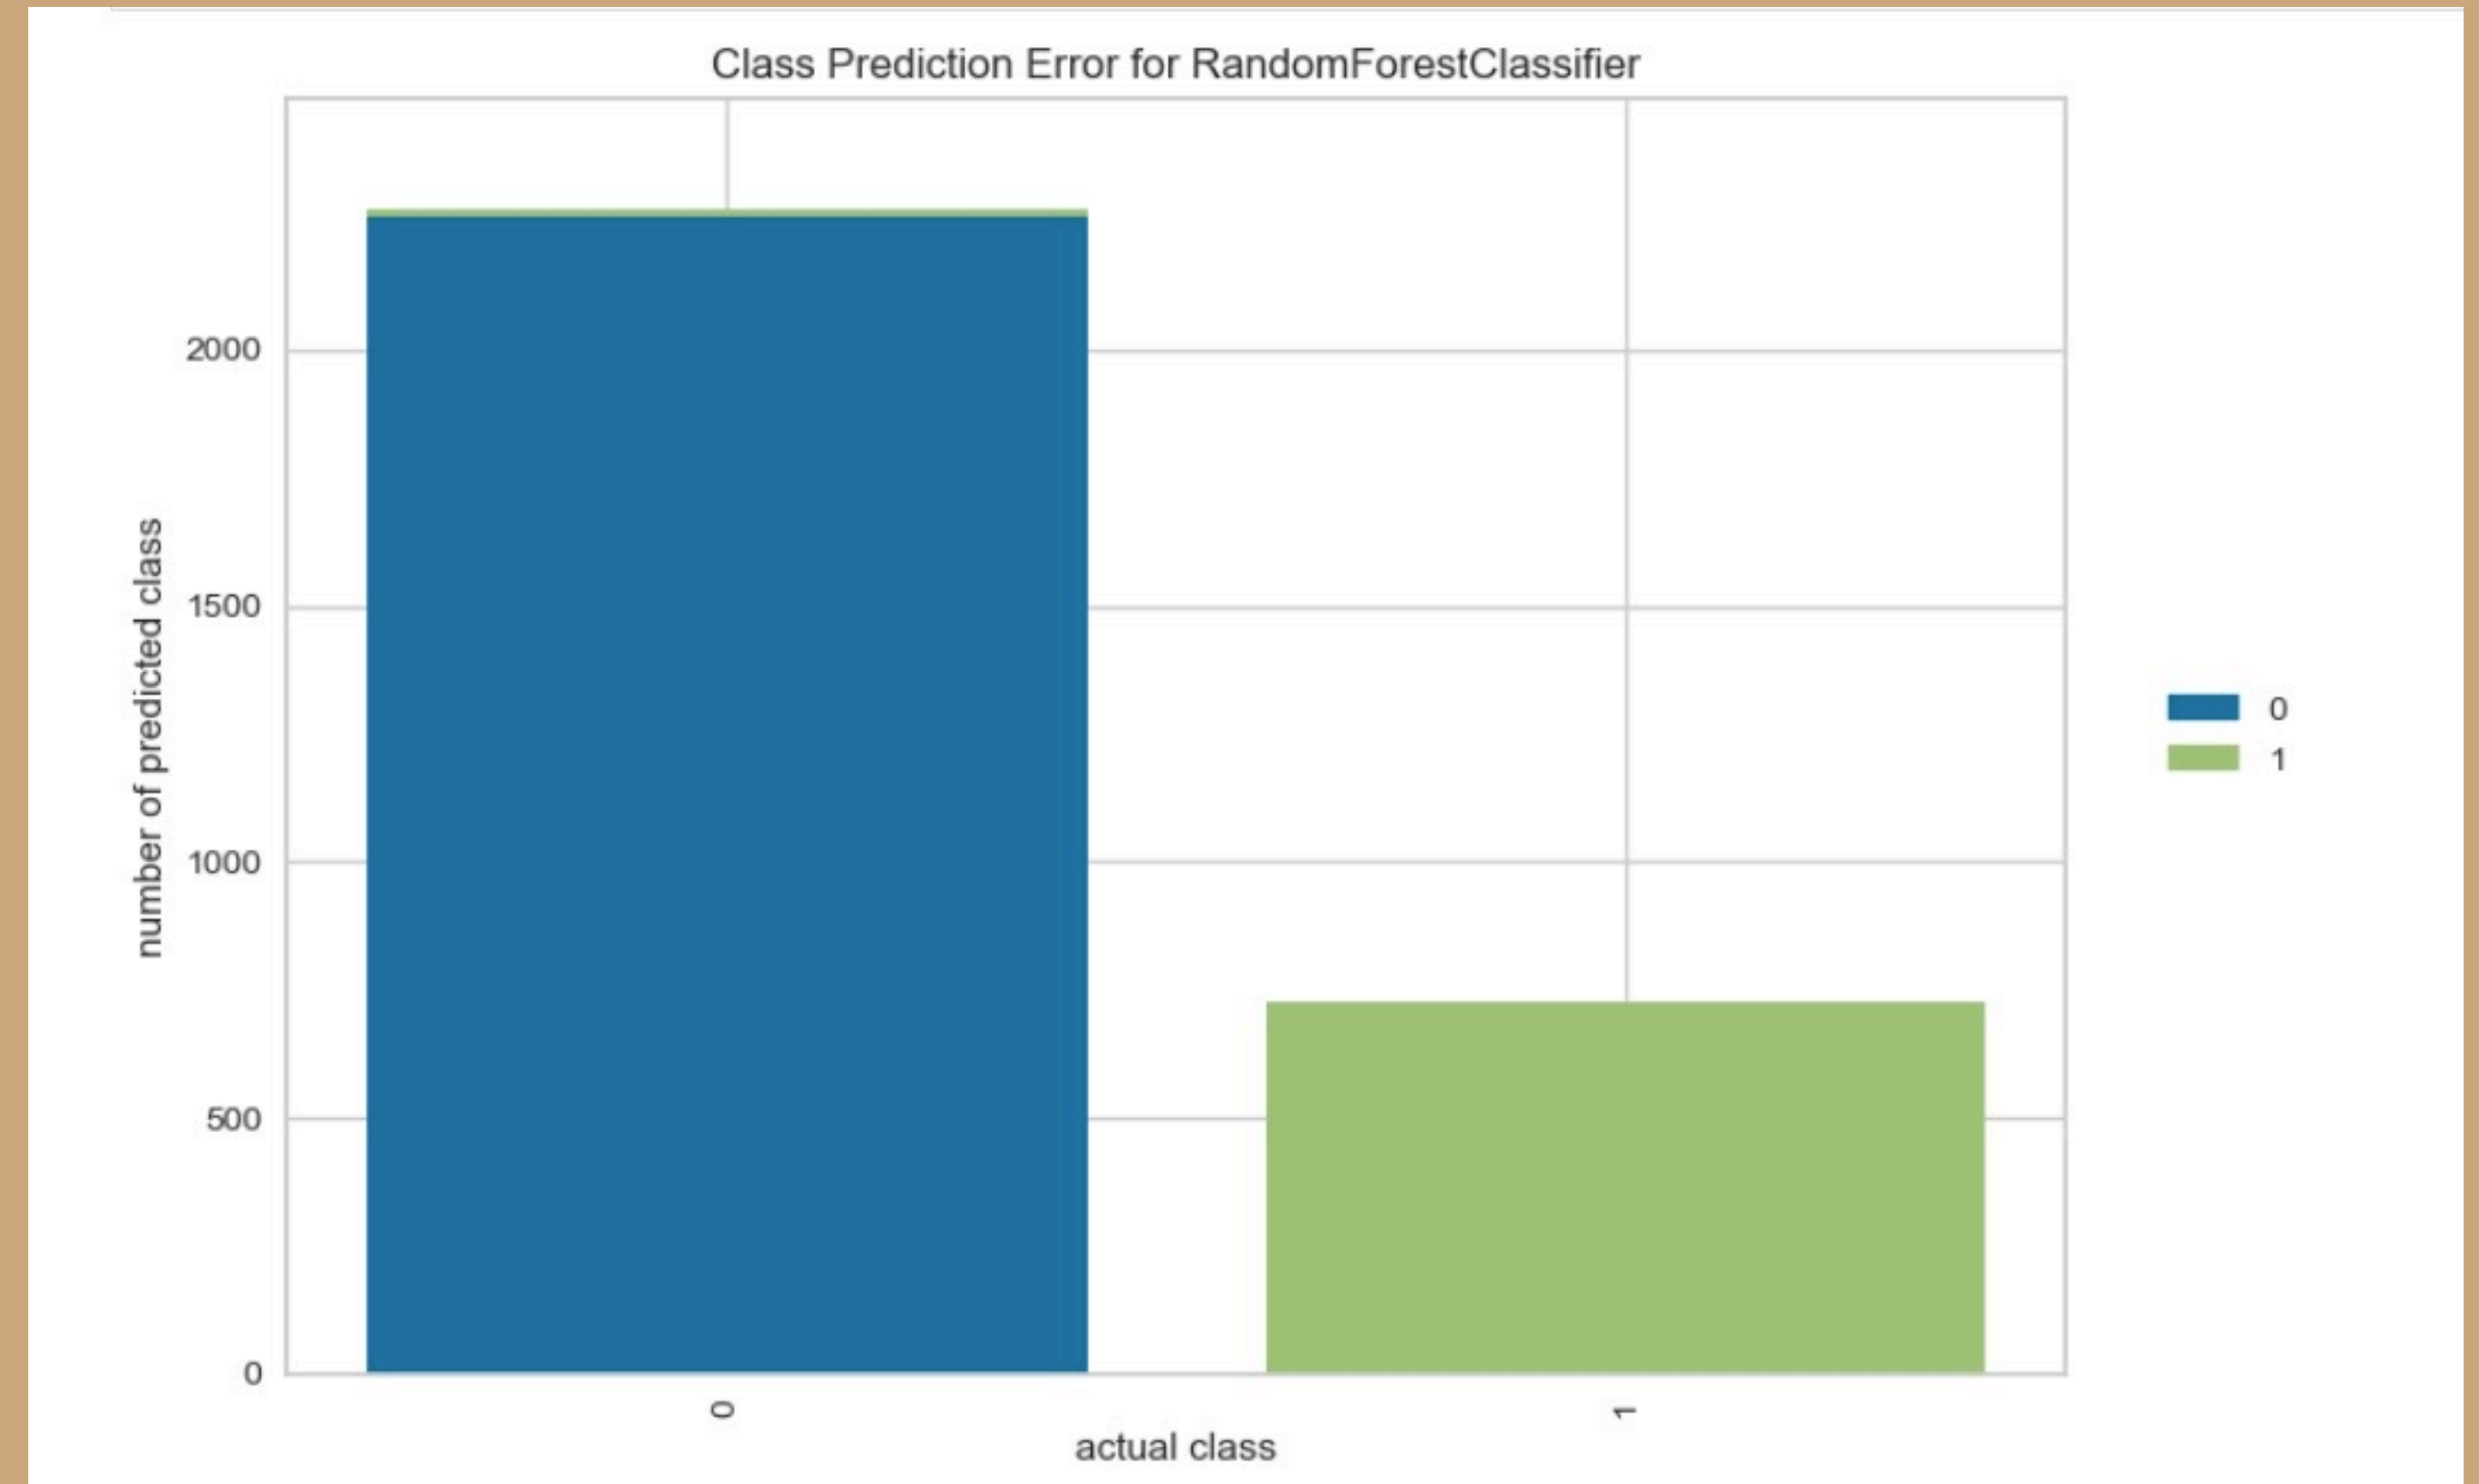
[[2239 23]  
[ 38 700]]

	precision	<u>recall</u>	f1-score	support
0	0.98	0.99	0.99	2262
<u>1</u>	0.97	<u>0.95</u>	0.96	738
<u>accuracy</u>			<u>0.98</u>	3000
macro avg	0.98	0.97	0.97	3000
weighted avg	0.98	0.98	0.98	3000

Train Set

[[9098 68]  
[ 137 2696]]

	precision	<u>recall</u>	f1-score	support
0	0.99	0.99	0.99	9166
<u>1</u>	0.98	<u>0.95</u>	0.96	2833
<u>accuracy</u>			<u>0.98</u>	11999
macro avg	0.98	0.97	0.98	11999
weighted avg	0.98	0.98	0.98	11999






# Model Building

## XGBoost

<u>Test_Set</u> [[2257 5] [ 15 723]]				
	precision	<u>recall</u>	f1-score	support
0	0.99	1.00	1.00	2262
<u>1</u>	0.99	<u>0.98</u>	0.99	738
<u>accuracy</u>			<u>0.99</u>	3000
macro avg	0.99	0.99	0.99	3000
weighted avg	0.99	0.99	0.99	3000
<u>Train Set</u> [[9166 0] [ 0 2833]]				
	precision	<u>recall</u>	f1-score	support
0	1.00	1.00	1.00	9166
<u>1</u>	1.00	<u>1.00</u>	1.00	2833
<u>accuracy</u>			<u>1.00</u>	11999
macro avg	1.00	1.00	1.00	11999
weighted avg	1.00	1.00	1.00	11999






# Model Building


## ANN

```
47/47 [=====] - 0s 1ms/step
[[1098   45]
 [  28  329]]
```

	precision	<u>recall</u>	f1-score	support
0	0.98	0.96	0.97	1143
<u>1</u>	0.88	<u>0.92</u>	0.90	357
<u>accuracy</u>			<u>0.95</u>	1500
macro avg	0.93	0.94	0.93	1500
weighted avg	0.95	0.95	0.95	1500



# Model Comparing



ROC\_AUC

# Model Deployment

## Select Employee's Attributes

Satisfaction

0,09

Last Evaluation

0,36

Number of Projects

1

7

Average Monthly Hours

1

310

Experience

1

10

Work Accident

Yes

Promotion

Yes

Department

sales

Salary

low

## Employee Churn Prediction App



Whether your employees will continue to work with you or not ? Let's See !

Please fill the attributes on the left hand side to make run the model properly.

	satisfaction	last_eval	num_of_projects	avg_monthly_hours	work_accident	experience	promotion
0	0.0900	0.3600	1	1	1	1	1

Check the features you selected from the table above. If correct, press the Predict button.

Predict

Your employee will \*\* leave.\*\*

**Any question/feedback?**

**Thank you!**